INF1060: Introduction to Operating Systems and Data Communication

Data Communication:

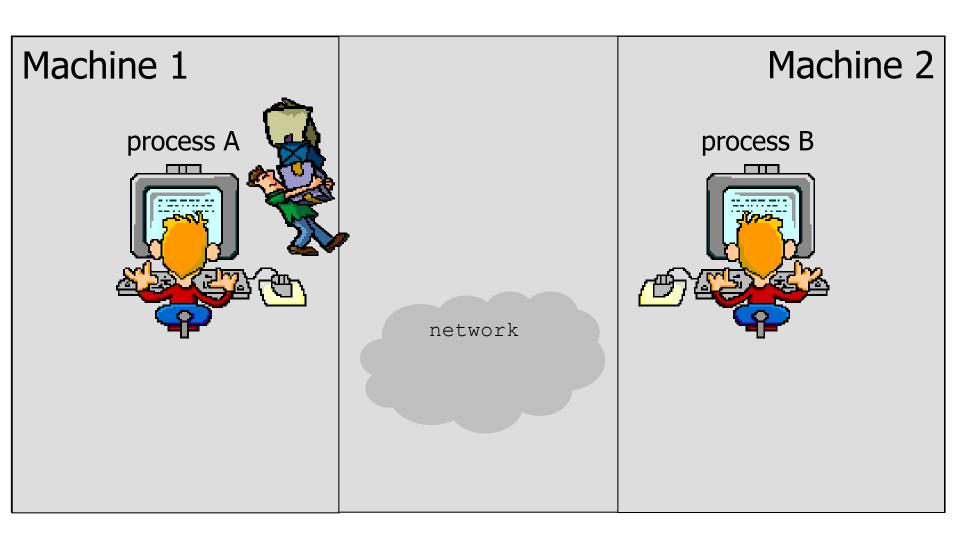
Introduction to Berkeley Sockets

Michael Welzl

(revised by Hans Petter Taugbøl Kragset 2015)

(adapted from lectures by Pål Halvorsen, Carsten Griwodz & Olav Lysne)

Big Picture



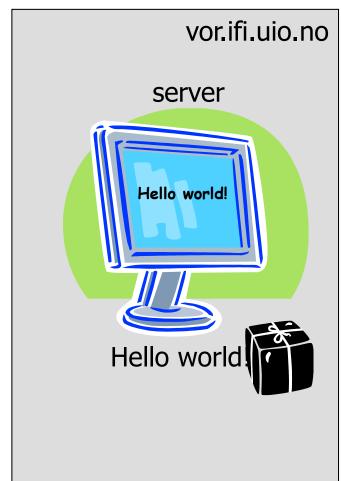
Goal

- Introduce socket API
- We will write two programs
 - A "client" and a "server"
- Each will run on one machine
 - For instance "vor.ifi.uio.no" (129.240.65.59) (maybe outdated IP)
- They will work as follows (see next slide)
 - The client sends the text "Hello world!" to the server
 - The server writes the received text on the screen
 - The server sends the received text back to the client and quits
 - The client writes the received text onto the screen and quits

What we want







What we want

Client

```
int main()
  char buf[13];
  /* Send data */
 write(sd, "Hello world!", 12);
  /* Read data from the socket */
  read(sd, buf, 12);
  /* Add a string termination sign,
     and write to the screen. */
 buf[12] = '\0';
 printf("%s\n", buf);
}
```

Server

```
int main()
 char buf[13];
 /* read data from the sd and
     write it to the screen */
  read(sd, buf, 12);
 buf[12] = '\0';
 printf("%s\n", buf );
 /* send data back over the connection */
 write(sd, buf, 12);
```

Read & Write

- Same functions used for files etc.
- The call read (sd, buffer, n);
 - Reads up to n characters
 - From socket sd
 - Stores them in the character array buffer

- The call write (sd, buffer, n);
 - Writes up to n characters
 - From character array buffer
 - To the socket sd

Alternatives to Read & Write

- The call recv(sd, buffer, n, flags);
 - - Used to control the behavior of the function
 - Several flags can be specified at once with bitwise or operaions
 - MSG_DONTWAIT | MSG_MORE
- The call send(sd, buffer, n, flags);
 - Flags, same as above
- Several similar functions like ...to/from, ...msq

Creation of a connection

- One side must be the active one
 - Take the initiative in creating the connection
 - This side is called the *client*
- The other side must be passive
 - It is prepared for accepting connections
 - Waits for someone else to take initiative
 - This side is called the server
- From now: server is a process, not a machine

Special for the server side

In case of TCP

- One socket on the server side is dedicated to waiting for a connection
- For each client that takes the initiative, a separate socket on the server side is created
- This is useful for all servers that must be able to serve several clients concurrently (web servers, mail servers, ...)

To do – in the code

Client

```
<Necessary includes>
int main()
  char buf[13];
  <Declare some more data structures>
  <Create a socket called "sd">
  <Identify the server that you want to contact>
  <Connect to the server>
  /* Send data */
  write(sd, "Hello world!", 12);
  /* Read data from the socket */
  read(sd, buf, 12);
  /* Add a string termination sign,
     and write to the screen. */
  buf[121 = '\0']
  printf("%s\n", buf);
  <Closing code>
```

Server

```
<Necessary includes>
int main()
  char buf[13];
  <Declare some more data structures>
  <Create a socket called "request-sd">
  <Define how the client can connect>
  <Wait for a connection, and create a new socket "sd"</p>
   for that connection>
  /* read data from the sd and
     write it to the screen */
  read(sd, buf, 12);
  buf[12] = '\0';
  printf("%s\n", buf );
  /* send data back over the connection */
  write(sd, buf, 12);
  <Closing code>
```

To do-list

Client:

- Declare some data structures
- Create a socket
- Identify the server
- Connect to the server

Server:

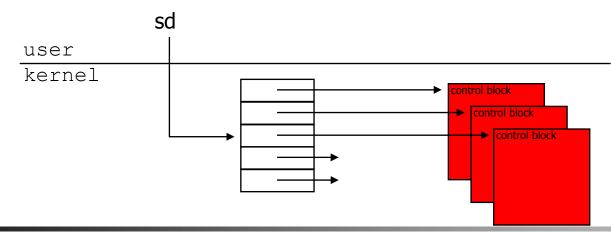
- Declare some data structures
- Create a socket
- Define how the client can connect
- Wait for a connection
- Handle the connection

<Create a socket>

Client

Server

 Call to the function socket() creates a transport control block (hidden in kernel), and returns a reference to it (integer used as index)



More about the socket call

```
sd = socket(int domain, int type, int protocol);
```

- PF_INET, SOCK_STREAM and IPPROTO_TCP are constants that are defined in the included files
- The use of the constants that we use here (as above) creates a TCP socket
- Many other possibilities exist
 - Domain: PF_UNIX, PF_INET, PF_INET6, ...
 - Type: SOCK STREAM, SOCK DGRAM, ...
 - Protocol: IPPROTO TCP, IPPROTO UDP, ...
- protocol can be NULL, OS choses apropriate proocol
 (use with care!)

How to identify clients to accept, and servers to contact?

- Machine??
 - by its IP address (e.g., 129.240.65.59)
- Application/service/program??
 - by (IP address and) port number
 - standard applications have own, "well-known" port numbers
 - SSH: 22
 - Mail: 25
 - Web: 80
 - Look in /etc/services for more

Address structure

struct sockaddr_in :

```
    sin_family address family used (defined through a macro)
    sin_port 16-bit transport protocol port number
    sin_addr 32-bit IP address defined as a new structure in_addr having one s_addr element only
    sin zero padding (to have an equal size as a sockaddr struct)
```

- Defines IP address and port number in a way the Berkeley socket API needs it
- man 7 ip

Address structure - handout

 Fill address type ("family"), address and port number into the structure

- a constant indicating that Internet protocols will be used
- INADDR ANY
 - a constant meaning any (Internet) address
 - in this context: any own Internet address

Address structure

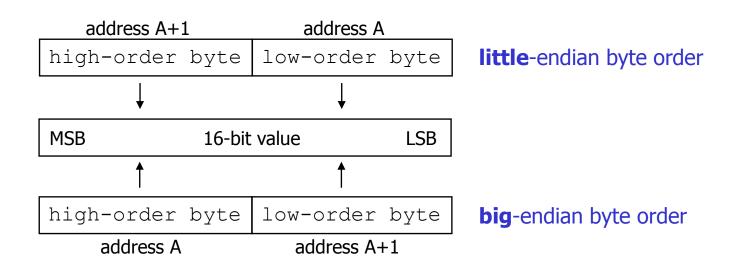
Client

```
/* declaration */
struct sockaddr in serveraddr;
/* clear the structure */
memset(&serveraddr, 0,
      sizeof(struct sockaddr in));
/* This will be an address of the
* Internet family */
serveraddr.sin family = AF INET;
/* Add the server address - vor */
inet pton (AF INET,
          "129.240.65.59",
          &serveraddr.sin addr);
/* Add the port number */
serveraddr.sin port = htons(2009);
```

Server

Byte Order

- Different machines may have different representation of multi-byte values
- Consider a 16-bit integer: made up of 2 bytes



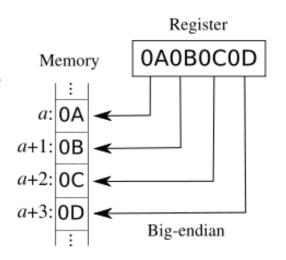
Byte Order: Storing 32-bit 0x0A0B0C0D

Assuming 8-bit (one byte) atomic elements...

...big endian:

- the most significant byte (MSB), 0x0A, is stored on the *lowest* memory address
- the least significant byte (LSB), 0x0D, is stored on the highest memory address

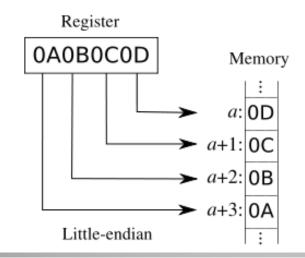
increasing memory addresses								
	0x0A	0x <mark>0B</mark>	0x0C	0x0D				



... little endian:

- 0x0A is stored on the highest memory address
- 0x0D is stored on the <u>lowest</u> memory address

increasing memory addresses								
	0x0D	0x0C	0x <mark>0B</mark>	0x0A				



Byte Order: IP address example

- IPv4 host address: represents a 32-bit address
 - written on paper ("dotted decimal notation"): 129.240.71.213
 - binary in bits: 10000001 11110000 01000111 10001011
 - hexadecimal in bytes: 0x81 0xf0 0x47 0x8b
- Big-endian ("normal" left to right):
 - one 4 byte int on PowerPC, POWER, Sparc,



- Little-endian:
 - one 4 byte int on x86, StrongARM, XScale, ...: (0x8b47f081)



Network byte order:

0x**81f0478b**

Byte Order: Translation

- Byte order translation makes communication over several platforms possible
- htons() / htonl()
 - host-to-network short / long
 - translate a 16 / 32-bit integer value to network format
- ntohs() / ntohl()
 - network-to-host short/long
 - translate a 16 / 32-bit integer value to host format
- Little-endian (x86 etc.): ntohl(0x81f0478b) == 0x8b47f081
- Big-endian (PowerPC etc.): ntohl(0x81f0478b) == 0x81f0478b

Presentation and Numeric Address Formats

- The network...
 - ...does not interpret the "dotted decimal notation" presentation format
 - ...needs a *numeric* binary format in network byte order
- inet pton()
 - translate the text string to a numeric binary format needed by the address structure
- inet ntop()
 - translate the (numeric binary) network address structure to a text string

How far have we gotten now?

Client

```
<Necessary includes>
int main()
  char buf[13];
<Declare some more data structures>
 . <Create a socket called "sd">
 <Identify the server that you want to contact>
Connect to the server>
  /* Send data */
  write(sd, "Hello world!", 12);
  /* Read data from the socket */
  read(sd, buf, 12);
  /* Add a string termination sign,
      and write to the screen. */
  buf[12] = '\0';
  printf("%s\n", buf);
<Closing code>
```

Server

```
<Necessary includes>
int main()
  char buf[13];
< Declare some more data structures>
<Create a socket called "request-sd">
<Define how the client can connect>
<Wait for a connection, and create a new socket "sd"</p>
   for that connection>
  /* read data from the sd and
      write it to the screen */
  read(sd, buf, 12);
  buf[12] = '\0';
  printf("%s\n", buf );
  /* send data back over the connection */
  write(sd, buf, 12);
<Closing code>
```

Bind and Listen

- bind(int sfd, struct sockaddr *a, socklen t al)
 - A machine can have several addresses (several network cards, loopback, ...)
 - Tells the socket on the server side which local protocol (i.e., *IP address* and *port number*) to listen to

- listen(int sfd, int backlog)
 - Prepares the server for listening to connect requests, and initializes a queue for connect requests
 - The second parameter (often SOMAXCONN) defines how long the queue(s) should be

Connect and Accept

- connect(int sfd, struct sockaddr *serv_a, socklen_t al)
 - connects client socket to a server that is specified in the address structure
 - a three-way handshake is initiated for TCP
 - possible errors
 - ETIMEDOUT no response (after several tries) and timer expired
 - ECONNREFUSED server not running or not allowed to connect
 - EHOSTUNREACH HOST not reachable
 - ENETUNREACH NET not reachable
- sd = accept(int sfd, struct sockaddr *a, socklen t *al)
 - take the first connect request from the connect request queue
 - wait for the connect request to arrive if the queue is empty
 - returns a **new socket** that the server can use to communicate with the client
 - a (clientaddr) contains information about the client
 - al must be initialized, so accept knows size of a

Binding, Listening, Accepting and Connecting

Client

Server

```
/* Bind the address to the socket */
bind(request sd,
     (struct sockaddr*) &serveraddr,
     sizeof(struct sockaddr in);
/* Activate listening on the socket */
listen(request sd, SOMAXCONN);
/* Create client addr struct*/
  clientaddrlen =
          sizeof(struct sockaddr in);
/* Wait for connection */
sd = accept(request sd,
          (struct sockaddr*) &clientaddr,
          &clientaddrlen);
```

Closing of Sockets

Client

/* Close the socket */ close(sd);

Server

```
/* Close both sockets */
close(sd);
close(request_sd);
```

- Note that the semantics of close depends
 - On the kind of protocol
 - Some possible extra settings
 - (similar for file descriptors used to operate on disk...)

All data that has not been read yet may be thrown away

Complete Client

Client

```
#include <netinet/in.h>
#include <sys/socket.h>
#include <netdb.h>
#include <stdio.h>
#include <string.h>
int main()
  /* Declarations */
  struct sockaddr in serveraddr;
  int sd:
  char buf[13];
  /* Create socket */
  sd = socket(PF INET,
                SOCK STREAM,
                IPPROTO TCP);
  /* Clear address structure */
  memset(&serveraddr, 0,
        sizeof(struct sockaddr in));
  /* Add address family */
  serveraddr.sin family = AF INET;
```

Client ctd.

```
/* Add IP address of vor.ifi.uio.no */
inet pton (AF INET, "129.240.65.59",
          &serveraddr.sin addr);
/* Add the port number */
serveraddr.sin port = htons(2009);
/* Connect */
connect (sd.
        (struct sockaddr*) & serveraddr,
        sizeof(struct sockaddr in));
/* Send data */
write(sd, "Hello world!", 12 );
/* Read data */
read(sd, buf, 12);
/* add string end sign, write to screen*/
buf[12] = '\0';
printf("%s\n", buf);
/* Close socket */
close(sd);
```

Complete Server

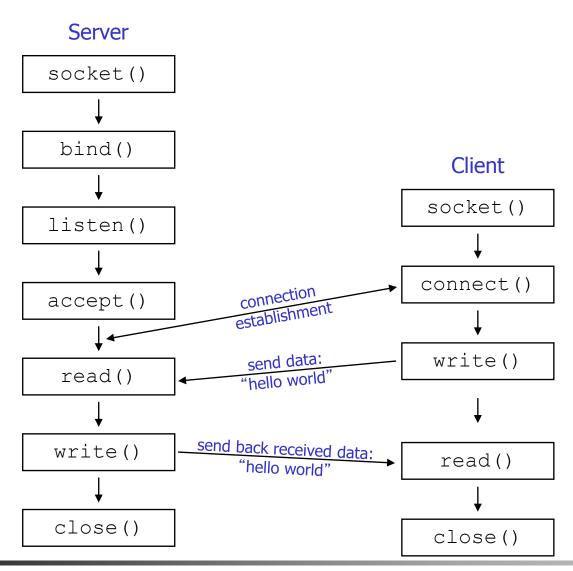
Server

```
#include <netinet/in.h>
#include <sys/socket.h>
#include <netdb.h>
#include <stdio.h>
#include <string.h>
int main()
  /* Declarations */
  struct sockaddr in serveraddr;
  struct sockaddr in clientaddr;
  int clientaddrlen:
  int request sd, sd;
  char buf[13];
  /* Create socket */
  request sd = socket(PF INET,
                        SOCK STREAM,
                        IPPROTO TCP);
  /* Fill in the address structure */
  memset(&serveraddr, 0,
        sizeof(struct sockaddr in));
  serveraddr.sin family = AF INET;
  serveraddr.sin addr.s addr = INADDR ANY;
  serveraddr.sin port = htons(2009);
```

Server ctd.

```
/* Bind address to socket */
bind(request sd,
     (struct sockaddr*) &serveraddr,
     sizeof(struct sockaddr in));
/* Activate connect request queue */
listen(request sd, SOMAXCONN);
/* Receive connection */
clientaddrlen =
    sizeof(struct sockaddr in);
sd = accept(request sd,
          (struct sockaddr*) &clientaddr,
          &clientaddrlen);
/* Read data from socket and write it */
read(sd, buf, 12);
buf[12] = '\0';
printf("%s\n", buf);
/* Send data back over connection */
write(sd, buf, 12);
/*Close sockets */
close(sd); close(request sd);
```

Summary of Socket Functions for our Elementary TCP Client-Server



Get the code yourself

- Similar examples can be found on the INF1060 Github page
- Also other examples!!

Complete Server – improvements?

Server

```
int main()
  /* Declarations */
  /* Create socket */
  request sd = socket(...);
  /* Fill in the address structure */
  /* Bind address to socket */
 bind(...);
  /* Activate connect request queue */
  listen(...);
```

Server ctd.

```
/* Receive connection */
sd = accept(...);
/* Process the request*/
/*Close sockets */
close(sd);
close(request sd);
```

Iterative servers?

Iterative Servers

Server

```
int main()
  /* Declarations */
  /* Create socket */
  request sd = socket(...);
  /* Fill in the address structure */
  /* Bind address to socket */
 bind(...);
  /* Activate connect request queue */
  listen(...);
```

Server ctd.

```
for (;;) {
  /* Receive connection */
  sd = accept(...);
  /* Process the request*/
  /*Close sockets */
  close(sd);
close(request sd);
```

Concurrent servers?

Concurrent Iterative Servers

Server

```
. . .
int main()
  /* Declarations */
 pid t pid;
  /* Create socket */
  request sd = socket(...);
  /* Fill in the address structure */
  /* Bind address to socket */
 bind(...);
  /* Activate connect request queue */
  listen(...);
```

Server ctd.

```
for (;;) {
 /* Receive connection */
  sd = accept(...);
  if ((pid = fork()) == 0) {
     close(request sd);
    /* Process the request*/
    /*Close sockets */
    close(sd);
   exit(0)
  /*Close sockets */
  close(sd);
close(request sd);
```

Select

- Problems with these examples:
 - iterative: cannot serve more than one socket at once
 - concurrent: overhead (a process per socket)
- Solution: functions that tell you when a socket becomes available (select, poll)
- int select(int nfds, fd_set *restrict readfds, fd_set
 *restrict writefds,fd_set *restrict errorfds, struct
 timeval *restrict timeout)
 - check whether fd's (sockets) from the nfds set are available for reading (readfds), writing (writefds), or have exceptional conditions pending (errorfds)
 - Null argument: don't check. Timeout = time limit for check (Null = block).
 - result is given by changing readfds / writefds / errorfds

Select usage and macros

Select usage

 Declare and initialize fd_set; add relevant sockets to fd_set;
 give select a copy of fd_set for every operation of interest (read/write/exceptional); loop through copies to take action

Preparing fd_set is done with some macros

- FD_CLR(fd, &fdset)
 - removes the socket descriptor fd from the socket descriptor set fdset
- FD ISSET(fd, &fdset)
 - returns nonzero if socket descriptor fd is a member of fdset; else 0
- FD_SET(fd, &fdset)
 - adds socket descriptor fd to fdset
- FD_ZERO(&fdset)
 - initializes fdset to 0, representing the empty set
- FD SETSIZE max. number of FDs; use this as the first parameter for select



University of Oslo INF1060

Complete Select-based Server

Test with e.g. two clients!

Server

```
#include <netinet/in.h>
#include <sys/socket.h>
#include <netdb.h>
#include <stdio.h>
#include <string.h>
#include <time.h>
int main()
  /* Declarations */
  struct sockaddr in serveraddr;
  struct sockaddr in clientaddr;
  int clientaddrlen, i, rc;
  int request sd, sd[2], numsocks, maxsocks;
  char buf[13];
  fd set fds, readfds;
  struct timeval timeout;
  numsocks = 0; maxsocks = 2;
  timeout.tv sec = 20;
  timeout.tv usec = 0;
  /* Create socket */
  request sd = socket(PF INET,
                        SOCK STREAM,
                        IPPROTO TCP);
```

Server ctd.

```
/* Fill in the address structure */
memset(&serveraddr, 0,
      sizeof(struct sockaddr in));
serveraddr.sin family = AF INET;
serveraddr.sin addr.s addr = INADDR ANY;
serveraddr.sin port = htons(2009);
/* Bind address to socket */
bind(request sd,
     (struct sockaddr*) & serveraddr,
     sizeof(struct sockaddr in));
/* Activate connect request queue */
listen(request sd, SOMAXCONN);
/* Initialize fd set */
FD ZERO(&fds);
FD SET(request sd, &fds);
```

Complete Select-based Server ctd.

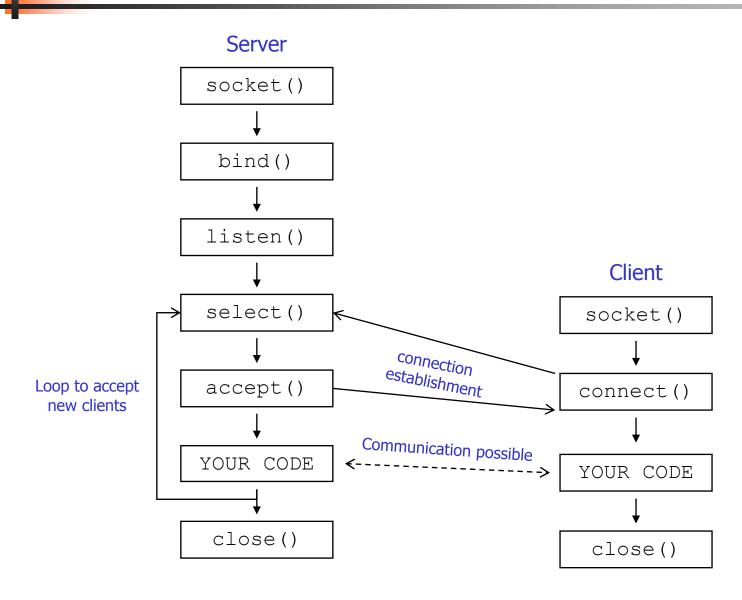
Server ctd.

```
for (;;) {
  readfds=fds;
  rc=select(FD SETSIZE, &readfds, NULL,
                  NULL, &timeout);
  /* Something went wrong */
  if (rc<0)
    return -1;
  /* Nothing happened, select continued */
  if (rc==0) {
    printf("Timeout!\n");
    for(i=0; i<numsocks; i++) {</pre>
      /* Send a response */
      write(sd[i], "Server ACK!",11);
      /* Close sockets */
      close(sd[i]);
      FD CLR(sd[i], &fds);
    return 0;
```

Server ctd.

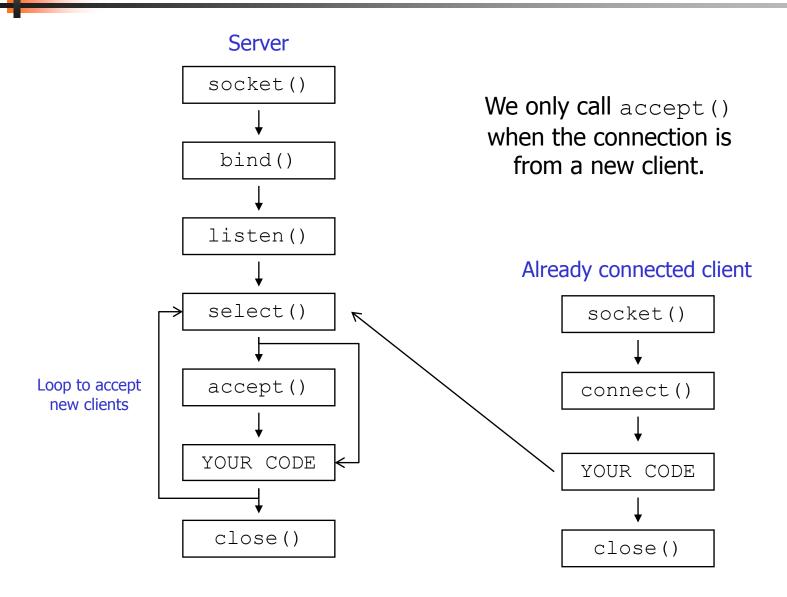
```
for (i = 0; i < FD SETSIZE; i++)
    if(FD ISSET (i, &readfds)) {
      if(i == request sock) {
        /* new connection request */
        if(numsocks < maxsocks) {</pre>
          sd[numsocks] = accept(request sock,
            (struct sockaddr *) &clientaddr,
            (socklen t *)&clientaddrlen);
          FD SET(sd[numsocks], &fds);
          numsocks++;
        } else {
          printf("Ran out of socket space.\n");
          return -1;
      } else {
        /* data arrived on an existing socket */
        read(i, buf,12);
        buf[12] = ' \ 0';
        printf("From socket %d: %s\n",i,buf);
close(request sock);
```

Outline for select based iterative server



42

Outline for select based iterative server



Summary

 We have implemented a short program where two processes communicate over a network

Next: the magic of how data is sent...

Literature

- "Berkeley UNIX System Calls and Interprocess Communication", Lawrence Besaw, University of Wisconsin
- Many books:
 - Kurose/Ross, "Computer Networking: A Top-Down Approach Featuring the Internet", 2nd ed., Addison-Wesley
 - Andrew Tanenbaum, "Computer Networks", 4th ed., Prentice Hall
 - W. Richard Stevens, "Unix Network Programming Networking APIs: Sockets and XTI", volume 1, 2nd ed., Prentice Hall